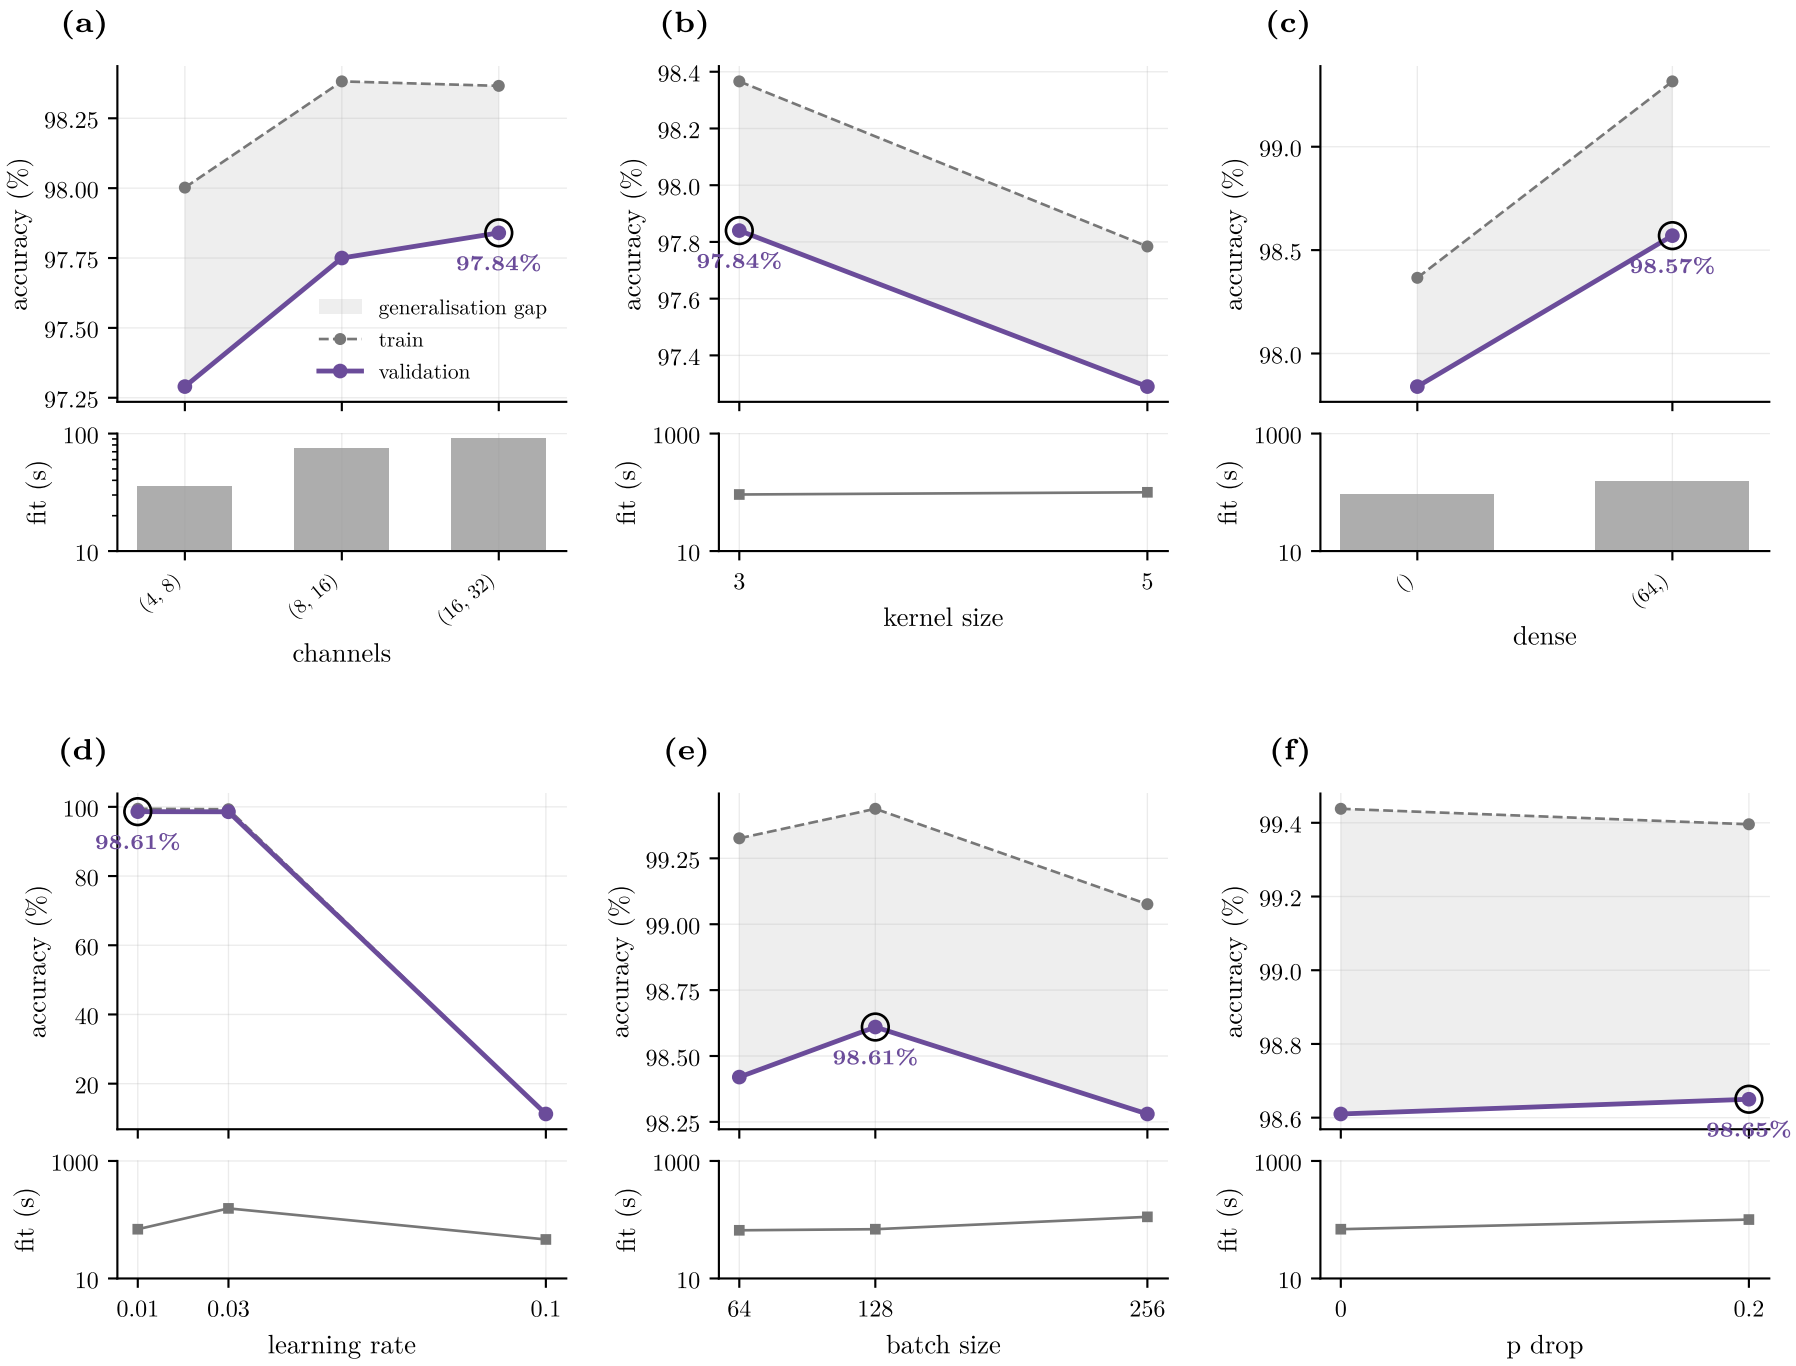
CNN: one-factor sweeps on the validation split (circle = selected value; lower strip = fit time)
(a)
(b)
(c)
(d)
(e)
(f)
accuracy (%)
fit (s)
generalisation gap
train
validation
97.84%
channels
(4, 8)
(8, 16)
(16, 32)
kernel size
dense
()
(64,)
learning rate
batch size
p drop
98.61%
98.57%
98.65%